

Physics 40 Lab Manual

Michael Mulhearn

October 21, 2021

Contents

1	Installation of Scientific Python	3
1.1	Introduction	3
1.2	Installing Miniconda3	3
1.2.1	Installing under Windows	4
1.2.2	Installing on a Chromebook	4
1.2.3	Installing Miniconda3 under Linux or macOS	4
1.3	Installing the Physics 40 Conda Environment	5
1.4	Starting a Jupyter notebook	5
1.5	Submitting your assignment	8
2	Binary Numbers	9
2.1	Introduction	9
2.2	Preparation	9
2.3	Binary Representation of Integers	10
2.4	Binary Representation of Real Numbers	12
2.5	Other Data Types	14
3	Sequences and Series	16
3.1	Introduction	16
3.2	Preparation	16
3.3	Fibonacci Sequence	17
3.4	Arithmetic Series	17
3.5	Geometric Series	18
3.6	Refinements	18
3.7	Fibonacci Integer Right Triangles	19
4	The Quadratic Equation and Prime Numbers	20
4.1	Introduction	20
4.2	Preparation	20
4.3	Quadratic Formula	21
4.4	Prime Numbers	22
4.5	The Lucky Number of Euler	23
5	Arrays, Plotting and Chaos	24
5.1	Introduction	24
5.2	Preparation	24
5.3	Plotting with Scientific Python	25
5.4	The Logistics Map	27

6	Differentiation and Projectile Motion	30
6.1	Introduction	30
6.2	Preparation	30
6.3	Numerical Differentiation	31
6.4	Projectile Motion	32
6.5	Projectile Motion with Drag	34
7	Pendulum Motion	35
7.1	Introduction	35
7.2	Preparation	35
7.3	Pendulum Motion	36
7.4	Small Angle Approximation	36
7.5	The Failure of the Euler Method	38
7.6	The Verlet Method	39
8	Orbital Motion	41
8.1	Introduction	41
8.2	Orbital Motion	41
9	Integration	42
9.1	Introduction	42
9.2	Integration	42
9.3	Gaussian Distribution	42
9.4	Elliptic Function	42

Chapter 1

Installation of Scientific Python

1.1 Introduction

In this lab, you will install the software which we will be using in phy40. This is an assignment, and will be graded. You should submit a text file containing a log of all the steps you took to install the software on your computer. Make this log as specific as possible, an entry might be:

Downloaded windows installer from:

https://repo.anaconda.com/miniconda/Miniconda3-latest-Windows-x86_64.exe

Keeping this log will also make it easier for you to get help if you have problems.

If you run into problems, do some research on a web search tool (Google, for example) to become better informed and to see if you can overcome the problem on your own before asking for help. This is an important technique in getting help with technical problems that will serve you well even outside of this class. You will find it more easy to get useful technical help, from the sort of people most capable of offering it, when it is clear from your question that you are informed and have already tried all of the obvious things. If you are still stuck after trying to solve the problem for yourself, then contact your TA or instructor with specific technical details about what is failing, and include your installation log.

If you do find a problem with these instructions or manage to overcome a technical problem yourself, make sure to note it in your log and inform your TA, in case it is helpful for other students.

1.2 Installing Miniconda3

We will be using Miniconda3 based on Python 3.7 for data analysis using Jupyter notebooks. Miniconda is a lightweight package which we can use to install all of the remaining analysis software we will need in a consistent manner across all different operating systems.

Determine which OS type and version you have on the desktop or laptop computer that you will be using for your coursework. The software here will work under Windows, Linux, or macOS. It should also work on all Chromebooks released since 2019, and some earlier Chromebooks. You should also check whether you have a 32-bit or 64-bit OS (you can find instructions for how to determine this for your particular OS version with a Google search.) Most desktop or laptop computers built in the last ten years are 64-bit.

If you are using Linux or macOS, then from within a terminal type:

```
echo $SHELL
```

to determine the shell you are using (typically "bash" these days). Record all of this information in your installation log file.

Once you have determined your OS type and version, follow the instructions below appropriate to your operating system.

1.2.1 Installing under Windows

If you have already installed a version of conda (e.g. Anaconda or Miniconda) then you do not need to re-install it. Instead, find the the Anaconda Prompt in the Application menu and run it.

If you need to install Miniconda3, then download and run the appropriate installer from:

<https://docs.conda.io/en/latest/miniconda.html#>

If prompted, you should choose to:

- Accept the license / terms of use.
- Install for just the current user, not all users.

Once installed, check that you can run the "Anaconda Prompt". From the prompt, check that you can run:

```
conda --version
```

and note the output in your installation log. Then proceed to Section 1.3.

1.2.2 Installing on a Chromebook

You will need to activate Linux on your Chromebook, according to the instructions here:

<https://www.codecademy.com/articles/programming-locally-on-chromebook>

Then follow the instructions for installing under Linux. If your Chromebook predates 2019 and does not support Linux, contact your instructor for alternative arrangements.

1.2.3 Installing Miniconda3 under Linux or macOS

If you believe you already already have a version of conda installed (such as miniconda or ananconda), check by running

```
conda --version
```

If you see something like:

```
conda 4.9.2
```

as output (even if the version is different) then you do indeed already have conda installed, with the base environment activated, and you can skip ahead to Section 1.3. If instead you get a message like:

```
conda: command not found
```

then the easiest solution is to simply proceed with these instructions.

To install Miniconda, download the appropriate installer for your OS here:

```
https://docs.conda.io/en/latest/miniconda.html\#
```

For macOS, you can choose between a "package" or "bash" version. I find it easier to follow the bash version, but the package version will work too. I recommend you make the following choices if prompted:

- Accept the license / terms of use.
- Do not install for all users, but just one the current user.
- Do allow the installer to issue "conda init".

During the installation, take note of the install location in your log.

After installation with these settings, conda will automatically activate the "base" conda environment. If this annoys you, as it does me, or interferes with other software you are using, you can turn off this aggressive behavior with:

```
conda config --set auto_activate_base false
```

Confirm that you have successfully installed conda by typing

```
conda --version
```

Record the output in your installation log, and proceed to Section 1.3.

1.3 Installing the Physics 40 Conda Environment

Make sure your conda is fully up to date with:

```
conda update conda
```

Then follow the prompts, e.g. selecting "y" as needed to update any out-of-date packages.

We'll be using a conda environment specifically for phy40 to avoid conflicts with any other projects on your computer, and to ensure that we all have the same software installed. To create our environment:

```
conda create -n phy40 python=3.9 numpy scipy matplotlib ipython jupyter
```

1.4 Starting a Jupyter notebook

This course will make extensive use of the Jupyter Notebook interface to Scientific Python, which is well suited to academic work (including independent research) because it combines code with output in digestible chunks. Even when the end product is a polished piece of software, much of the initial code development can be done in the interactive session that Jupyter Notebooks provide.

To activate the phy40 environment type:

```
conda activate phy40
```

When you are done with Phy 40 for the day you can deactivate this environment (later) with:

```
conda deactivate
```

Launch jupyter notebook with:

```
jupyter notebook
```

```

mulhearn@vonnegut: ~/lab1
File Edit View Search Terminal Help
mulhearn@vonnegut:~$ conda activate phy40
(phy40) mulhearn@vonnegut:~$ mkdir lab1
(phy40) mulhearn@vonnegut:~$ cd lab1
(phy40) mulhearn@vonnegut:~/lab1$ jupyter notebook
[I 16:27:03.601 NotebookApp] Serving notebooks from local directory: /home/mulhearn/lab1
[I 16:27:03.601 NotebookApp] Jupyter Notebook 6.4.3 is running at:
[I 16:27:03.601 NotebookApp] http://localhost:8888/?token=919c1f7fbfdc2e78e8f36ff1ae1b082ebd85b86dc3443e84
[I 16:27:03.601 NotebookApp] or http://127.0.0.1:8888/?token=919c1f7fbfdc2e78e8f36ff1ae1b082ebd85b86dc3443e84
[I 16:27:03.601 NotebookApp] Use Control-C to stop this server and shut down all kernels (twice to skip confirmation).
[C 16:27:03.640 NotebookApp]

To access the notebook, open this file in a browser:
file:///home/mulhearn/.local/share/jupyter/runtime/nbserver-22257-open.html
Or copy and paste one of these URLs:
http://localhost:8888/?token=919c1f7fbfdc2e78e8f36ff1ae1b082ebd85b86dc3443e84
or http://127.0.0.1:8888/?token=919c1f7fbfdc2e78e8f36ff1ae1b082ebd85b86dc3443e84

```

Figure 1.1: Example starting Jupyter Notebook from the Linux command line. In Windows, you will need to open the Anaconda Prompt instead of a terminal.

```

In [1]: # Lab 1 - Installation of Scientific Python
        # Michael Mulhearn (Working Independently)
        %pylab inline

        Populating the interactive namespace from numpy
        and matplotlib

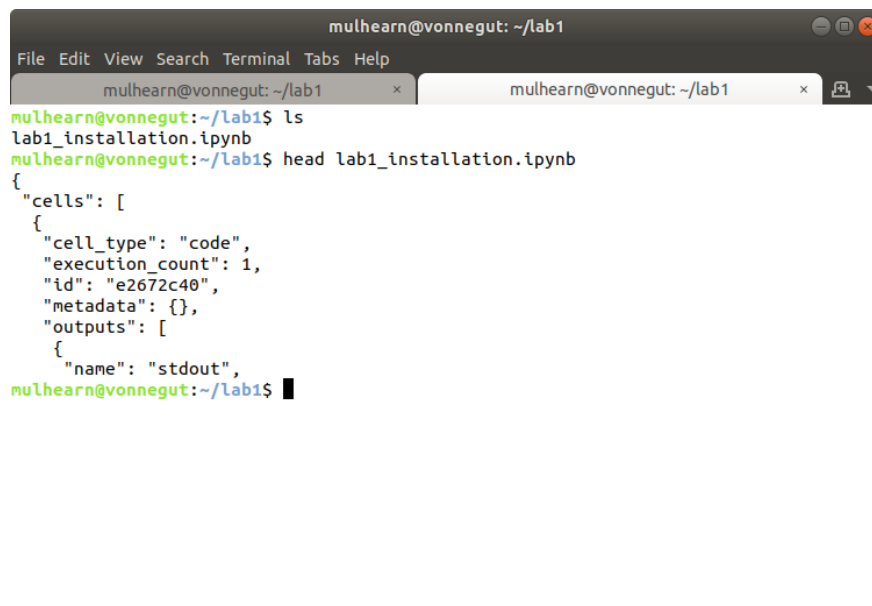
In [2]: #1.1
        print("Hello World")

        Hello World

In [ ]:

```

Figure 1.2: The Hello World example Jupyter Notebook.



```

mulhearn@vonnegut: ~/lab1
File Edit View Search Terminal Tabs Help
mulhearn@vonnegut: ~/lab1 x mulhearn@vonnegut: ~/lab1 x
mulhearn@vonnegut:~/lab1$ ls
lab1_installation.ipynb
mulhearn@vonnegut:~/lab1$ head lab1_installation.ipynb
{
  "cells": [
    {
      "cell_type": "code",
      "execution_count": 1,
      "id": "e2672c40",
      "metadata": {},
      "outputs": [
        {
          "name": "stdout",
mulhearn@vonnegut:~/lab1$ █

```

Figure 1.3: Example showing the saved Jupyter notebook. Notice that notebook file (ipynb) is not human readable on its own: it requires the Jupyter software to render it in a human readable form.

This should start the Jupyter Notebook server and open a client in your web browser. An example starting a Jupyter Notebook from Linux is shown in Fig. 1.1.

You should create one Jupyter Notebook per lab assignment, by choosing the New (Python 3) option in your client. Change the name of your notebook to something that clearly identifies the lab. Start each lab with comments (starting with “#” symbol) indicating the title of the lab, then your name followed by your lab partners. See the first cell of Fig. 1.2 for an example. This first cell is also a good place to issue the ipython “magic function”:

```
%pylab inline
```

which will setup the notebook for inline plots and load the numpy and matplotlib libraries for you.

Each assignment will consist of a number of steps, clearly numbered like this one, your first step:

△ **Jupyter Notebook Exercise 1.1:** Print “hello world” using the python print command.

To keep your notebook clear, label cells (such as this one) with a comment for the assignment step number, as in the second cell of Fig. 1.2. You only need to label one cell if the assignment is fulfilled across several cells.

Jupyter Notebook checkpoints your work automatically. You should be able to see your notebook saved in the working directory where you started, as in Fig. 1.3. Notice that while the notebook file is ASCII text, it is not a human readable format. The Jupyter software is needed to render the notebook in a human readable way. To make your grader’s life easier, you will be submitting PDF versions of your notebook, once all of the tasks are completed and the output is visible. There are several ways to make a PDF file from your notebook, but the most reliable is to use the “Print Preview” option to view the notebook as a PDF file within your browser, then use the print feature of your browser to print the page as a PDF file. Try this now, and make sure you can create a legible PDF file, but do not submit it to the course site, as you still have more to do. Always keep

your python notebook file (ipynb) even after you submit the assignment. If you have problems, you can reproduce a PDF file from the notebook file, but it is tedious to reproduce your notebook from PDF. If you have problems producing the PDF file, you can submit the “ipynb” file as a temporary work-around, but work with your TA to sort out the problem as quickly as possible.

1.5 Submitting your assignment

Before submitting, take some time to clean up your assignments to remove anything superfluous and place the exercises in the correct order. You can also add comments as needed to make your work clear. You can use the Cell → All Output → Clear and Cell → Run All commands to make sure that all your output is up to date with the cell source.

When you are satisfied with your work, print the PDF file as described earlier and submit **both** the PDF file and notebook file to the course website.

Chapter 2

Binary Numbers

2.1 Introduction

At the heart of numerical analysis, naturally, you will find numbers. In this lab, we will explore the basic data types in Python, with particular emphasis on the computer representation of integers and real numbers. All modern programming languages do an admirable job of hiding the limitations of the computer representations of these mathematical concepts. In this chapter, we will deliberately explore their limitations.

2.2 Preparation

This lab will rely on the material from Section 1.2.2 of the Scientific Python Lecture notes.

△ **Jupyter Notebook Exercise 2.1:** Enter the following code into a cell and check the output:

```
a = 121
print(type(a))
print(a)
```

Next, in the same cell, add a line at the end setting a to a real value, `a = 1.34`, and print the type and value again. Check the output. Next, set a to Boolean value, `a = False`, and print the type and value yet again.

In many languages, such as C and C++, variables are strictly typed: you would have to decide at the start whether you want a to be of integer, float, or boolean type, and then you would not be able to change to a different type later. Python variables are references to objects, which means they only point to memory locations that contain objects with all of the data and functionality associated with that object. When you write $a = 121$ it is interpreted as “set variable a to point to a location in memory that contains a class of type integer with the value 121”.

△ **Jupyter Notebook Exercise 2.2:** Enter the following code into a cell and check the output:

```
a = 12
b = a
b = 5
print(a)
print(b)
```

Why is the output “12, 5” instead of “5, 5”? When you write `b = a`, the variable `b` points to the same Integer class that `a` points to. So when you write “`b = 5`” why doesn’t the value of `a` change as well? This code snippet shows that it does not. The reason is that Integers are *immutable* objects in python... their values cannot be changed. So when you write `b = 5` it is interpreted as “variable `b` points to a (new) Integer with value 5.” The variable `a` continues to point to the Integer with value 12. The “is” operator used like this:

```
print(a is b)
```

tells you if `a` references the same object as `b`. Add to compiles of this line to your snippet in order to clarify the situation. One call should return True and the other False.

△ **Jupyter Notebook Exercise 2.3:** If I set a variable `a` to an integer value and I set `b` to the same integer value, do `a` and `b` refer to the same object, or to two different objects with the same value? Write a snippet of code (three lines) to find out.

In this lab, it will be convenient to know how to calculate the absolute value of a number, which you can do with the python `abs` function:

```
a = -1.5
b = abs(a)
print(b)
```

2.3 Binary Representation of Integers

Computer hardware is based on digital logic: the electrical voltage of a signal is either high or low, which correspond to a mathematical zero or one. A digital clock is used to ensure that signals are only sampled at particular times, when they are guaranteed not to be in transition from a zero to one or vice versa. The Arithmetic-Logic Unit (ALU) uses digital logic gates (such as AND or OR) to perform calculations. For example, it is possible to build an adder that uses only NAND gates.

Because digital signals have only two states (zero and one), the most natural way to represent numbers in a computer is using the base two, which we call binary. In the familiar base ten, we have ten digits (from 0 to 9) and the place value increases by a factor of 10 with each digit moving toward the left. In binary, we have only two digits (0 and 1) and the place value increase by factors of two. A single digit in binary is referred to as a “bit”. You add columns quickly when counting in binary: zero(0), one(1), two(10), three(11), four(100), five(101), and so on. For efficient operations, computers often group eight bits together to form a byte. Digital values are therefore commonly represented in hexadecimal (base 16) where two digits of hexadecimal describes one byte. See Table 2.1, which you can produce for yourself in Python like this:

```
print("dec: hex: bin:")
for d in range(16):
    print("{0:<2d}    0x{0:01x}    0b{0:04b}".format(d))
```

It is conventional to prepend binary numbers with “0b” and hexadecimal with “0x” otherwise we wouldn’t know whether a “10” represents ten, sixteen, or two! Notice that the largest number that can be written with n bits is $2^n - 1$.

The mathematical notion of an integer can be naturally implemented by computer hardware. Although integers are represented in binary in the hardware, modern compilers and languages generally print them to screen as decimal by default. One caveat is that computers do not have

Table 2.1: The numbers 0 to 15 in decimal, hexadecimal, and binary.

dec:	hex:	bin:	dec:	hex:	bin:
0	0x0	0b0000	8	0x8	0b1000
1	0x1	0b0001	9	0x9	0b1001
2	0x2	0b0010	10	0xa	0b1010
3	0x3	0b0011	11	0xb	0b1011
4	0x4	0b0100	12	0xc	0b1100
5	0x5	0b0101	13	0xd	0b1101
6	0x6	0b0110	14	0xe	0b1110
7	0x7	0b0111	15	0xf	0b1111

an unlimited number of bits. Many computer languages use 64-bit integers, which means that only the integer values from 0 to 18446744073709551615 can be represented:

```
x = 2**64-1
print(x)
```

For signed integers, one bit is used to indicate the sign (positive or negative) and so a 64-bit signed integer can represent integer values from -9223372036854775808 to 9223372036854775807. As long as an integer value is within the range covered by the integer type, the integer value can be perfectly represented.

Python uses arbitrary sized integers: it simply adds more bits as needed to represent any number. For an extremely large number, you will eventually reach practical limitations on the amount of memory and processing time available in the computer, which will limit how large of an integer can be calculated.

△ **Jupyter Notebook Exercise 2.4:** See for yourself just how huge integers can be in python by entering:

```
x = 2**8000
print(x)
```

and checking the output.

△ **Jupyter Notebook Exercise 2.5:** Print the integer 64206 in decimal, hexadecimal, and binary. Hint: just reuse the carefully formatted print statement from the example above.

△ **Jupyter Notebook Exercise 2.6:** Suppose you are tasked with rewriting the firmware for a distant satellite which has just lost one line from an eight-bit digital communications bus due to radiation damage. You now have only seven working bits! What is the maximum sized unsigned integer which you could write on this degraded seven-bit bus? What range of signed integers could you write?

△ **Jupyter Notebook Exercise 2.7:** This is a paper and pencil exercise. You can do it on paper and pencil and submit a scanned PDF, or you can just record you steps as comments in a jupyter notebook cell. Let's consider a four-bit signed integer, so zero is 0000 and one is 0001. Suppose the upper bit is reserved for sign, so 1XXX is a negative number. An obvious choice for representing

-1 would be 1001, but there is a better choice. Consider that:

$$(-1) + 1 = 0$$

Well, if we simply ignore the last carry bit (5th bit):

$$1111 + 1 = 10000 = 0000$$

So if we define 1111 as -1, we can treat addition with negative numbers exactly the same as adding ordinary numbers. Find the representation for -2 such that

$$-2 + 2 = 10000 = 0000$$

then show that:

$$-1 + -1 = -2.$$

Python does not use this trick, but many other languages do.

2.4 Binary Representation of Real Numbers

Representing real numbers presents much more of challenge. There are an uncountably infinite number of real numbers between any two distinct rational numbers, but a computer has only finite memory and therefore a finite number of states. It is impossible for computers to exactly represent every real number. Instead, computers represent real numbers with an approximate floating point representation much like we use for scientific notation,

$$x = m \times B^n$$

where the significand m is a real number with a finite number of significant figures and the exponent n is an integer. The base B is ten for scientific notation but typically two in a floating point representation. The exponent n is typically chosen so that there is only one digit before the decimal point in the base B , e.g. 3.173×10^{-8} for scientific notation.

The limited precision of the discriminant can lead to challenges when using floating point numbers. The floating point precision is specified by the parameter ϵ (epsilon) which is the difference between one and the next highest number larger than one that can be represented. For scientific notation with four significant digits, $\epsilon = 0.001$, because we cannot represent anything between 1.000×10^0 and 1.001×10^0 with only four significant figures.

△ **Jupyter Notebook Exercise 2.8:** Determine the Python floating point ϵ by running

```
import sys
print(sys.float_info.epsilon)
```

△ **Jupyter Notebook Exercise 2.9:** Determine the Python floating point ϵ for yourself by running:

```
eps = 1.0
while eps + 1 > 1:
    eps = eps / 2
eps = eps * 2
print(eps)
```

Here the while loop continues running the indented code until the condition $\epsilon + 1 = \epsilon$ is met.

△ **Jupyter Notebook Exercise 2.10:** Python uses the IEEE 754 double-precision floating-point format. This format uses 64-bits overall, with 52 bits reserved for the significant. The standard uses a clever trick to save one bit, by requiring that the leading bit of the significant m is one, and defining 53 significant figures using 52 bits:

$$m = 1.m_1m_2m_3\dots m_{52}$$

Here each m_i represents an individual bit (0 or 1). Predict the parameter ϵ and compare with the above.

△ **Jupyter Notebook Exercise 2.11:** It seems like ϵ should be small enough to simply ignore it in most cases, but in fact it shows up quite clearly if you apply strict equality to floating point quantities. To see the problem, run this code, checking if $\sin(\pi)$ is zero:

```
x = np.sin(np.pi)
print(x)
print(x==0)
```

The strict equality condition $x == 0$ is not met because of limited floating point precision. Instead of strict equality, check that x is near zero with a condition like:

$$|x| < 10\epsilon$$

where ϵ is the machine precision and the factor of 10 is a conservative factor to allow for round-off errors that might be a bit larger than the best possible precision ϵ . Add such a condition to the code above and show that this looser definition of equality now holds. In general, when using approximations (like floating point numbers) you can only check things to within the accuracy of the approximation!

△ **Jupyter Notebook Exercise 2.12:** Here is another case where floating point precision joins the chat uninvited:

```
x = 0.1
y = x+x+x
print(y == 0.3)
```

Devise an alternative to $y == 0.3$ that properly accounts for floating point precision.

△ **Jupyter Notebook Exercise 2.13:** Personally, I prefer my zero's to look like zero. When floating point limitations are making them look non-zero, I like to clean them up with rounding, like this:

```
x = np.sin(2*np.pi)
print(x)
x = np.around(x,15)
print(x)
```

Yeck! What is -0 ?! IEEE 754 defines two zeros -0 and 0 . -0 is used to indicate that 0 was reached by rounding a negative number. This is so that $1/-0$ can be interpreted as $-\infty$ and $1/0$ as $+\infty$. If you just want to make this go away, add “ $+0$ ”:

```
x = np.around(x,15)+0
```

Show that this works.

△ **Jupyter Notebook Exercise 2.14:** Consider the following code:

```
a = 5;
b = 1.2343E-17;
sum = 0
sum += 5;
for i in range(1000000):
    sum = sum + b
print(sum)
```

Is there any problem here? If there is, fix it by changing only the *order* of the lines of code.

2.5 Other Data Types

Integers and floating point numbers are the real work horses of computational physics. We'll add numpy arrays in a future lab. This section will briefly introduce the remaining types.

Python includes strings as a basic type:

```
s = "hello world"
s = s + " (it's been a strange few years)"
print(s)
print(type(s))
print(s[6], s[4], s[6])
print(type(s[0]))
```

Strings are *immutable* objects that contain textual data. If you have used other languages, you might expect `s[0]` to be a “character” but in Python it is a string of length one. There is no built-in character type.

Python includes complex numbers as a basic type:

```
z = 1 + 2j
print(type(z))
print(z.real)
print(z.imag)
print(z.conjugate())
```

This is our first example of an object oriented programming (OOP) class interface. To compute the complex conjugate of z , an ordinary function would need to be passed z as an argument, or else it would not know which complex value to use for the computation. But `z.conjugate()` is a *method* of the class `complex`. The method is tied to the instance of complex number z by the “.” and has access to all of the data it needs from z . Similarly, the `real` and `imag` are member data of class `complex`: they are the integers that contain the real and imaginary parts of z . Objects play a central role in Python, but in a refreshingly understated and reserved manner. It is enough for now to understand that `z.conjugate()` is much like a function that already has z as a parameter, and `z.imag` and `r.real` are just ways to access the data contained in z .

△ **Jupyter Notebook Exercise 2.15:** Define a complex number with value of $1+i$ and multiply it by its complex conjugate. Show that the resulting complex number has zero for its imaginary part.

Python provides Lists as a native container of python objects. We'll make much more use of numpy arrays, which are better suited to numerical analysis, but Python Lists occasionally play a role for various bookkeeping tasks:

```
L = ["hello", 1, 2, 3+2j, 3.45, "green"]
print(L[0])
print(L[1])
print(L[5])
L[0] = "goodbye"
print(L)
```

Here the List L contains a variety of (admittedly rather useless) objects. These objects can be referred to individually by their index. One place where lists really shine is in looping over a custom list of values:

```
for i in [1,5,10,50,100]:
    print(i)
```

△ **Jupyter Notebook Exercise 2.16:** Run the following code:

```
a = [1,2,3,4,5]
b = a
b[0] = 3
print(a[0])
print(a is b)
```

(Spits out coffee) “What the???!!” Lists are *mutable* which makes them fundamentally different from *immutable* integers. Here we assign b to point to the same object as a (a list) and then change an entry in that list. Even after the change, a and b point to the same object. This is only possible because the list object is mutable.

△ **Jupyter Notebook Exercise 2.17:** It's good to read the documentation, but it's a useful skill to figure things out for yourself too! Without looking up the documentation, write a snippet of code to determine for yourself if complex numbers are mutable (like lists) or immutable (like integers). We are able to change just a part of a list (like $a[0] = 3$). But can we change part of a complex number (like $z.\text{real}$). Try it and find out!

(It's OK if your code throws an error here, but you can also comment it out if you are the sort of person that can't possibly leave it alone)

Chapter 3

Sequences and Series

3.1 Introduction

In this lab, we will apply for loops to study sequences and series. If you already have programming experience, you can complete a challenge problem in lieu of some of the other problems: see the final problem for the details.

3.2 Preparation

This lab will rely on the material from Sections 1.2.1 to 1.2.4 of the Scientific Python Lecture notes. Most of the problems can be completed using a simple functions containing a single for loop, such as in this function:

```
def loop(n):  
    for i in range(n):  
        print(i)
```

To run the code in the function, you call the function, usually in a different cell:

```
loop(5)
```

△ **Jupyter Notebook Exercise 3.1:** Create a new function:

```
def mult(a,n):  
    # your code here ...
```

that prints the first n multiples of a. For example `mult(3,4)` should output:

```
3  
6  
9  
12
```

In future problems, we'll describe this output simply as 3, 6, 9, 12. We won't be picky about whitespace unless we discuss it explicitly. One way to complete this is to use the three arguments of `range(start,stop,step)`.

3.3 Fibonacci Sequence

The Fibonacci numbers are a sequence of numbers satisfying the recursion relationship:

$$F_{n+2} = F_n + F_{n+1}$$

with $F_0 = 0$ and $F_1 = 1$. The sequence is:

$$0, 1, 1, 2, 3, 5, 8, 13, 21, 34, \dots$$

This sequence can be generated numerically by an algorithm such as this one:

```

1  fa := 0 # set fa to 0
2  fb := 1 # set fb to 1
3  repeat n times:
4      fc := fa + fb
5      print fc to screen
6      fa := fb
7      fb := fc

```

Note that this is not python syntax. What is the importance of the last two lines? Would the algorithm work if we exchanged their order?

△ **Jupyter Notebook Exercise 3.2:** Use the algorithm described above to implement a new function `fib(n)` which prints the next n Fibonacci numbers after the initial 0 and 1.

3.4 Arithmetic Series

The finite arithmetic series

$$S_n = \sum_{k=0}^n (a + kd) = a + (a + d) + (a + 2d) + \dots + (a + nd)$$

sums to the average of the first and last terms times the number of terms:

$$S_n = (n + 1) \frac{a + (a + nd)}{2} \quad (3.1)$$

We will assume $a = d = 1$ and calculate this finite series numerically using the following function:

```

def arith(n):
    sum = 0
    for j in range(1, n+1):
        sum = sum + j
        #print("j: ", j, "\t sum: ", sum)
    return sum

```

Type in this function and see how it works by uncommenting the print statement (delete the `#` symbol that starts a comment) and calling it as `arith(5)`. The use of print statements in a loop like this or at each stage of a calculation is a simple, effective and classic debugging technique. You test your code with the print statements included, keeping n small so you don't fill your whole screen with output. Once your code is working, you comment out the unneeded print statements so that the interpreter ignores them and you no longer see the unneeded output. Why not

just delete them? You can, but experience shows that if you do, you will need the line again shortly!

△ **Jupyter Notebook Exercise 3.3:** Obtain the sum of the first n terms of arithmetic series with `sum = arith(n)` for three different values of n . Each time, show that sum returned by the function matches the expected sum.

3.5 Geometric Series

The geometric series

$$\sum_{k=0}^{\infty} ar^k = a + ar + ar^2 + ar^3 + \dots$$

converges for $|r| < 1$ to:

$$\frac{a}{1-r}. \quad (3.2)$$

We will demonstrate this numerically.

△ **Jupyter Notebook Exercise 3.4:** Implement a function `geom(a,r,n)` which calculates sum of the first n terms of the geometric series with k th term ar^k . Show that it agrees with Eqn. 3.2 for $a = 2$, $r = 0.5$ $n = 100$.

△ **Jupyter Notebook Exercise 3.5:** Call you geometric series function again for $a = 3$, $r = 0.8$ and $n = 100$. Compare with the expected output calculate within python and with pencil and paper. Do they agree exactly? If not, do they agree within the floating point precision?

△ **Jupyter Notebook Exercise 3.6:** Now compare your calculated sum with Eqn. 3.2 for $a = 1$, $r = -0.9$ $n = 100$. How is the agreement? Increase n and see what happens. Why do you suppose this series is slower to converge?

3.6 Refinements

There are a few refinements you can make to your code. Don't change your working code from previous examples! Instead, copy the previous version to a new cell and make your refinements there. You don't even need to change the name of the function, Python will happily overwrite the old function implementation when it reaches the cell with the new version. Make these code improvements:

△ **Jupyter Notebook Exercise 3.7:** (Optional) Improve your Fibonacci function so that prints the first n numbers including the initial two numbers "0" and "1". Make sure it works properly for $n = 0$, $n = 1$, $n = 2$, and so on.

△ **Jupyter Notebook Exercise 3.8:** (Optional) Extend the Arithmetic series function to include parameters a and d . Show that it works.

3.7 Fibonacci Integer Right Triangles

Starting with the number 5, every second Fibonacci number is the length of the hypotenuse of a right triangle with integer sides. The first two are:

$$5^2 = 3^2 + 4^2$$

and

$$13^2 = 5^2 + 12^2.$$

Furthermore, from the second triangle onward, the middle side is the sum of the lengths of the sides of the previous triangle, for example:

$$12 = 3 + 4 + 5.$$

△ **Jupyter Notebook Exercise 3.9:** (Optional Challenge) Use numerical methods to explicitly verify these properties for the first n Fibonacci integer right triangles.

If you would prefer, you may submit the Optional Challenge problem plus the problems from Section 3.5 to complete the assignment.

Chapter 4

The Quadratic Equation and Prime Numbers

4.1 Introduction

In this lab, we will make more extensive use of conditional statements to implement algorithms which solve the quadratic equation, identify prime numbers, and add fractions. An optional challenge problem, The Lucky Number of Euler, explores how the quadratic equation can generate prime numbers.

4.2 Preparation

This lab will rely on the material from Sections 1.2.1 to 1.2.4 of the Scientific Python Lecture notes. We'll now be making frequent use of conditional statements:

```
def compare(a,b):  
    if (a==b):  
        print("a equals b")  
    elif (a<b):  
        print ("a is less than b")  
    else:  
        print ("a is greater than b")
```

Notice that Python uses `==` for comparison. You will get a syntax error if you use `a=b` instead.

We'll also use of the modulo operator `%` extensively. The modulo operation $a\%b$ returns the remainder from the integer division a/b .

```
# is b a factor of a?  
def isfactor(a,b):  
    if (a%b == 0):  
        return True;  
    return False
```

Why does `a%b == 0` mean that `b` is a factor of `a`?

△ **Jupyter Notebook Exercise 4.1:** Consider this verbose code snippet:

```
for i in range(100):
    print("on index ", i)
```

Which prints the current index on every iteration. Use the modulo operator to modify the code so that it only prints the index every 10 iterations. This is a classic trick!

We will also be using while loops, which repeat a block of code until a condition is met:

```
count=0
while(count<10):
    print(count)
    count = count+1
```

4.3 Quadratic Formula

The Quadratic equation:

$$ax^2 + bx + c = 0$$

has solutions which are given by the quadratic formula

$$x = \frac{-b \pm \sqrt{b^2 - 4ac}}{2a} \quad (4.1)$$

The number of unique real solutions depends on the quantity in the square root, which is called the discriminant:

$$b^2 - 4ac$$

If this is positive there are two real solutions, if it is one there is one real solution, and if it is negative there are no real solutions. For now, let's assume that a, b, and c are all integers.

In this case, the solution to the quadratic equation can be calculated as follows:

```
1  D := b2 - 4ac
2  if (D=0):
3      calculate the single solution from quadratic
4      print single solution
5  if (D<0):
6      print no solutions
7  if (D>0):
8      calculate both solutions from quadratic formula
9      print both solution
```

A test case with one real solution is:

$$(x - 1)(x - 1) = x^2 - 2x + 1.$$

A test case with two real solution is:

$$(x - 1)(x + 1) = x^2 - 1.$$

A test case with zero real solutions is:

$$(x - i)(x + i) = x^2 + 1.$$

△ **Jupyter Notebook Exercise 4.2:** Implement a function `quad(a,b,c)` which reports the solutions to the quadratic equation and verify it with the test cases shown.

△ **Jupyter Notebook Exercise 4.3:** Calculate three more test cases with integer solutions and use them to test your function more thoroughly.

△ **Jupyter Notebook Exercise 4.4:** (Optional) The condition that the discriminant is exactly zero (`D == 0`) is problematic when applied to floating point numbers. (Why?) In example is for `a=.1` `b=.3` `c =0.225` which should have only one real solution. Test you function with this test case. Modify the conditionals in your function to account for floating point precision and test it.

4.4 Prime Numbers

A prime number is a number that has two and only two factors: itself and one. One is not prime, but two is. We can determine if a number a is prime as follows:

```

1  if (a<2):
2      return false
3  i := 2
4  while (i ≤ √a):
5      if (a%i==0):
6          return false
7      i := i + 1
8  return true

```

Why is there no need to check for factors larger than \sqrt{a} ?

△ **Jupyter Notebook Exercise 4.5:** Implement a function `isprime(a)` which returns `True` if the integer a is prime and `False` if not.

Suppose that next we want to find the first n prime numbers greater than or equal to a number A . We can simply check if A , $A + 1$, $A + 2$, and so on are prime until we find the first n . But we do not know how many numbers we will have to check before finding n that are prime. This is a case for a `while` loop.

△ **Jupyter Notebook Exercise 4.6:** Find the first n prime numbers greater than A using a while loop and your `isprime` function. Test it for $n = 10$ and $A = 0$, then for $A = 1000000000$. Try that with paper and pencil!

Notice how we broke this problem of finding primes into two parts: determining whether or not a number is prime or not, then testing and counting prime numbers. We thoroughly tested the first part before using it in the second. This is an essential approach to solving computational problems: break complicated tasks down into smaller tasks which can be tested separately.

△ **Jupyter Notebook Exercise 4.7:** Implement a function which computes the fraction

$$\frac{n}{d} = \frac{a}{b} + \frac{c}{d}$$

from integer inputs a, b, c and d and returns integers n and d . Returning $n > d$ is allowed, but n/d should be a simplified fraction with a greatest common factor of one. You can return two integers as a Tuple, like this:

```
# function which adds fractions
def addfrac(a,b,c,d):
    n=d=0
    #your code...
    return n, d

# calling function:
n,d =addfrac(1,2,1,3)
print("{0}/{1}".format(n,d))
```

For full credit, you must factorize (see what I did there?) this problem into two parts: your addfrac function should call a second function that does one well defined task.

4.5 The Lucky Number of Euler

Euler discovered that the formula:

$$k^2 + k + 41$$

produces prime numbers for $0 \leq k \leq 39$. Perhaps you can beat Euler at his own game!

Consider quadratics of the form:

$$k^2 + ak + b$$

For each integer value of a and b , there is a maximum number n such that the quadratic formula produces prime numbers for $0 \leq k < n$

△ **Jupyter Notebook Exercise 4.8:** (Optional) Find the the values of a and b which produce the largest number of prime numbers. Restrict yourself to $|a| \leq 1000$ and $|b| \leq 1000$.

If you do complete this optional problem, then nice work, hot shot, but remember that Euler found his without using a computer!

Chapter 5

Arrays, Plotting and Chaos

5.1 Introduction

This lab will introduce a fundamental element of scientific python, the numpy array, and use them to produce plots. We will also consider the non-linear logistics map and examine bifurcation in the approach to chaos.

If you can complete all of the plots in Section 5.4 including the optional plot and *without any instructor help* then you may omit the plots from the preceding sections.

5.2 Preparation

This lab will rely on the material from Sections 1.4.1 to 1.4.2 and 1.5.1 to 1.5.2 of the Scientific Python Lecture notes. This is the first lab that relies on inline plotting, so make sure you are starting your notebook with the “line magic”:

```
%pylab inline
```

This will load the numpy library as np, the matplotlib.pyplot library as plt, and setup the matplotlib backend to imbed plots in your notebook.

A Numpy array is a grid of values. Unlike Python lists, the elements of a numpy array all have the same data type, which makes them much more computationally efficient. Choices for the data type include the built-in python integer, float, and bool types. The numpy library provides a wide range of analysis tools that are mostly centered on the numpy array type.

Numpy arrays can be constructed easily from a Python list:

```
a = np.array([1.3, 7.2, 4.1, 0.0])
b = np.array([[1, 2], [3, 4]])
print(a)
print(b)
print(np.shape(a))
print(np.shape(b))
```

This is convenient when you have specific values you want to define by hand. Another possibility is to construct the numpy array by calling a function designed specifically for the purpose:

```
a = np.linspace(0, 1, 11)
print(a)
b = np.arange(0, 5, 1)
print(b)
```

Both `linspace` and `arange` allow you to specify the range of values you want, but with `linspace` you specify the number of points you want whereas with `arange` you specify the step size.

One of the great joys of using numpy arrays comes from the fact that most operators are applied elementwise automatically, without the need to explicitly write a for loop:

```
a = np.arange(0,5,1)
b = 10
print(a)
print(b)
print(a+b)
```

Notice how the value of b (10) is added to *every element* of a , without the need to explicitly loop over every element. Try modify the example code to multiply every element of a by b . Then try raising each element of a to the power of 2.

△ **Jupyter Notebook Exercise 5.1:** Use numpy arange and elementwise operations to implement a function `def powers(a, n)` which returns a numpy array containing the first powers of a from 1 to a^n . So for example `print(powers(2,4))` should output `[1 2 4 8 16]`

vskip 1cm

Numpy arrays can also be built up element by element using the append function:

```
a = np.array([])
print(a)
a = np.append(a, 1)
print(a)
a = np.append(a, 3)
print(a)
a = np.append(a, 4)
print(a)
```

This example creates an empty numpy array and then adds one element at a time.

△ **Jupyter Notebook Exercise 5.2:** Run the snippet above and observe the output.

△ **Jupyter Notebook Exercise 5.3:** Implement a function `def recur(n,x)` which returns an array containing the first n values of the recurrence relationship $x_{i+1} = 2x_i + 1$ starting from $x_0 = x$. Use the following algorithm:

```
1  Parameter x # starting value
2  Parameter n # number of iterations
3  Create empty array a
4  Repeat n times:
5      append x to array a
6      x := 2x+1
7  print a
```

Test your code for $x = 1$ and $n = 5$ and ensure that it produces the correct output: `[1. 3. 7. 15. 31.]`.

5.3 Plotting with Scientific Python

Basic plotting in Python requires two numpy arrays: one for the x coordinates and one for the y coordinates. Consider the following very simple plot:

```
x = np.array([0.0, 1.0, 2.0, 3.0, 4.0, 5.0])
y = np.array([0.3, 3.2, 5.8, 9.0, 12.4, 14.7])
plot(x,y,"bo")
```

Here, the “bo” options specifies blue circles. Now consider:

```
x = np.linspace(0, 1, 100)
y = np.sin(np.pi*x)
plt.plot(x,y,"r-")
```

Here the “r-” option specifies red line. Including 100 points (as done here) results produces a smooth looking curve.

Now promise me that you will never make another plot without labeling the x and y axes! Here’s another example will all the bells and whistles you need to make a professional looking plot:

```
UPPER = 2
LOWER = 0
tau = 2*np.pi
x = np.linspace(LOWER, UPPER, 100)
s = np.sin(tau*x)
c = np.cos(tau*x)
plt.plot(x,s,"b-",label="sin")
plt.plot(x,c,"r-",label="cos")
plt.xlabel("x")
plt.ylabel("y")
plt.title("Two Periods of a Sine and Cosine")
plt.legend(frameon=False)
plt.show()
```

Make sure you understand all of the features demonstrated here:

- Variables `UPPER` and `LOWER` located at the top of the snippet, allowing for easy adjustment of parameters that affect the plot.
- Use of `np.linspace` to define an array of x values, with plenty of them (100) to produce nice smooth curves.
- Creation of two different arrays of y values, one for sin and one for cos.
- Plotting the arrays of x and y values with `plt.plot` using the “-” option for a line and color blue(“b”) for sin and red(“r”) for cos.
- Defining appropriate axis labels with `plt.xlabel` and `plt.ylabel`.
- Adding a title with `plt.title`
- Creation of a legend using the `label` optional argument to `plt.plot` and the `plt.legend()` command. Removing the frame with option `frameon=False`

It is written so concisely and intuitively, you might not even notice what is going on with the line:

```
s = np.sin(tau*x)
```

Remember that x here is a numpy array of 100 elements. The `tau*x` multiplies every element of x by our value `tau`. The `np.sin(tau*x)` then takes the sine of each element. The resulting numpy array, also of 100 elements, is referenced by variable `s`. Each element of `s` contains $\sin(\tau x)$ for

the corresponding element of the array x . It takes some getting used to for programmers used to explicitly writing for loops for things like this, but ultimately, the fact that python handles so much of this bookkeeping for us is what makes it a very fun language to work with.

△ **Jupyter Notebook Exercise 5.4:** Plot the $\text{sinc}(x)$ function as a smooth line in the x range from -5 to 5. Add appropriate labels to each axes. Include a legend identifying the sinc function. For the line color, use any color other than red or blue.

5.4 The Logistics Map

The logistics map is the recurrence relation

$$p_{i+1} = r p_i (1 - p_i)$$

which calculates the value of p for step $i + 1$ from step i . The variable p can represent the ratio of a population to its maximum possible value, and each iteration (n) is a step in time (such as one year). Each year, the population increases due to birth and decreases due to starvation as the population approaches its maximum value (p near 1). The growth (or decline) of the population is controlled by the growth rate parameter r . We will only consider r in the range from $[0, 4]$ which keeps p in the range $[0, 1]$. This is a simple non-linear model which illustrates chaotic behavior.

△ **Jupyter Notebook Exercise 5.5:** Implement a function `def logmap(p, r)` which returns the next iteration (p_{i+1}) of the logistic map for parameter r and $p_i = p$. Test your code by showing that for $r = 3.0$ and $p_0 = 0.1$ the next five iterations are: 0.27, 0.5913, 0.725, 0.5981 and 0.7211.

△ **Jupyter Notebook Exercise 5.6:** Use your `logmap` function to build an array containing the first n entries in the logistics map starting from $p_0 = 0.1$. Check that the output is correct by comparing with the results from the previous exercise for $n = 6$.

△ **Jupyter Notebook Exercise 5.7:** Plot the time evolution of the logistics map for $r = 2.8$ for 40 iterations, starting from $p_0 = 0.1$ as in Fig. 5.1. To create a plot, you'll need an array containing the p values (these correspond to the y axis in the plot) which you can construct as in the previous exercise. But what about the x axis values? Since your array of p values contains $[p_0, p_1, p_2 \dots p_n]$ the corresponding array of indices is simply $[0, 1, 2 \dots 3]$ which you can construct using `np.arange`.

Your results from the previous exercise should reproduce Fig. 5.1. This shows that for $r = 2.8$ the logistics map converges to a value of about 0.64. But this non-linear recursion relationship does not always converge to a single value.

△ **Jupyter Notebook Exercise 5.8:** Plot the time evolution of the logistics map for $r = 3.2$ for 40 iterations, starting from $p_0 = 0.1$.

Notice that now the system oscillates between two values (near 0.5 and 0.8).

△ **Jupyter Notebook Exercise 5.9:** Plot the time evolution of the logistics map for $r = 3.5$ for 40 iterations, starting from $p_0 = 0.1$.




Figure 5.1: Convergence of the logistic map for $r = 2.8$




Figure 5.2: Long-term behavior of the logistics map as a function of parameter r .

Notice that now the system oscillates between four values. This is an example of bifurcation in the approach to chaos. To show this more clearly, we'd like to produce a plot as in Fig. 5.2. For each r value, this shows the p values from 100 iterations *after* the first 1000 iterations. This shows the *long term behavior* of the logistics map. We can clearly see that for $r = 2.8$ the p values converge to one single value and for $r = 3.2$ there are two values just as in the previous exercises. These bifurcations continue until the system becomes chaotic (oscillating between many different values) with occasional windows of stability.

To produce this plot for yourself, start by considering this snippet:

```
Nr = 5
r = np.linspace(2.6,4,Nr)
p = logmap(0.1,r)
print(p)
p = logmap(p,r)
print(p)
```

Here we create a numpy array of r values, and pass that to our `logmap` function instead of a single value. The operations within the function are applied elementwise to the array, and the result is that instead of a single p value, the call to `logmap(0.1,r)` returns an array of p values, one for each r value. This is just what we need to make the plot in Fig. 5.2.

△ **Jupyter Notebook Exercise 5.10:** Run (and understand!) the snippet and make a plot of p versus r . Understand that you are plotting p_2 as a function r ! Use the `"k,\"` format option (black pixels) when plotting. Comment out the print statements and increase Nr to 100.

△ **Jupyter Notebook Exercise 5.11:** Make a plot that is *almost* like that of Fig. 5.2 by plotting p_{1001} versus r . Hint: in the code above, instead of one call to `p = logmap(p, r)` use a for loop which calls this 1000 times.

You should start to see features of the Fig. 5.2 but you are only plotting one p value for each r . To see the bifurcations and chaos, you'll need to plot about 100 p values at each r value.

△ **Jupyter Notebook Exercise 5.12:** Reproduce the plot in Fig. 5.2. Hint: instead of plotting just p_{1001} as in the previous exercise, plot the next 100 values as well. Increase the number of r values plotted to 1000. Add appropriate labels to each axis.

△ **Jupyter Notebook Exercise 5.13:** (Optional) The bifurcation diagram which you have constructed exhibits self-similarity. If you zoom into an appropriate region of the diagram, you will find a diagram which looks quite similar to the original diagram. Produce a plot that demonstrates this self-similarity by zooming into a particular region.

Chapter 6

Differentiation and Projectile Motion

6.1 Introduction

In this lab we will apply numerical differentiation to elementary functions. We'll apply the Euler method to compute the trajectory of a projectile, and compare our results with the analytic solution. We'll also show that our numerical technique easily accommodates air resistance, a problem that has no analytic solution using elementary functions.

6.2 Preparation

Our code is going to get complicated enough that we will need to pay some attention to variable scope, as illustrated here:

```
i=1
j=2
k=3
def f(i,j):
    print(i,j,k)
f(i,j)
f(j,i)
```

Try to predict the output of this snippet before running it. The first three lines define integers i , j , and k . These have global scope, which means they can be accessed from anywhere. The function $f(i,j)$ has parameters i and j which have a scope limited to the function f . Even though they have the same name as the global variables i and j , they are independent quantities. Within the function f the variable i is the first parameter, and j is the second parameter. Because they have the same name, the global variables i and j are *shadowed* by the local parameters i and j . The global variable k is not shadowed.

In this lab, we will also be passing a function as an argument to another function, as in this simple example:

```
def show(f):
    print(f(2))
    print(f(3))

def f(i):
    return i**2

def g(i):
```

```

    return i**3

show(f)
show(g)

```

Here the `show` function takes another function `f` as an argument. Within the `show` function, the function `f` is called using parenthesis just like any other function, as in `f(2)`. We define two additional functions, `f` which returns the square of its argument, and `g` which returns the cube. When `show(f)` the output 4 and 9. When `show(g)` the output 8 and 27. Run the code as is, and also check what happens if you define `g(i)` to require a second argument as in `g(i,j)`.

6.3 Numerical Differentiation

In lecture we derived the right derivative (aka forward derivative) formula

$$f'(x) = \frac{f(x+h) - f(x)}{h} + \mathcal{O}(h)$$

for numerically determining the derivative of the function f . Remember we do not calculate the $\mathcal{O}(h)$ term, that indicates that the truncation error is of order h . We also derived the centered derivative formula:

$$f'(x) = \frac{f(x+h) - f(x-h)}{2h} + \mathcal{O}(h^2)$$

To evaluate a derivative using any of these formulas, we need to choose an appropriate value of h . If h is too large, the truncation error (the amount the estimated value differs from the actual value) will dominate. If h is too small, we will encounter problems with floating point precision.

△ **Jupyter Notebook Exercise 6.1:** Implement the right derivative formula as `right(f, x, h)` where f is the function to be evaluated, x is the location to evaluate the derivative, and h is the step size for the numerical integration. Check your code on several functions with known derivatives, like this:

```

def f(x): # derivative 0
    return 2
def g(x): # derivative 3
    return 3*x
def h(x): # derivative 4x
    return 2*x**2

print(right(f,1,0.01))
print(right(g,1,0.01))
print(right(h,1,0.01))

```

△ **Jupyter Notebook Exercise 6.2:** Implement the center derivative function as `center` and test it in the same manner as in the previous exercise for `right`.

△ **Jupyter Notebook Exercise 6.3:** Compare the performance of `right` and `center` like this:

```
def f(x): # derivative 6x**2
    return 2*x**3

print("right:", right(f,1,0.1), "center:", center(f,1,0.1))
print("right:", right(f,1,0.01), "center:", center(f,1,0.01))
print("right:", right(f,1,0.001), "center:", center(f,1,0.001))
```

Recall that the truncation error goes as h for the right derivative and as h^2 for the center derivative. Are these results consistent with that expectation?

From now on, we will use the center derivative function only due to its better performance. We can plot the derivative of a function like this:

```
def f(x):
    return 0.5*x**2

x = np.linspace(0,1,100)
y = center(f, x, 0.1)
plt.plot(x,y,"-b")
plt.xlabel("x")
plt.ylabel("y")
plt.show()
```

Notice how the argument x passed to the function `right(f,x,h)` and then to `f(x)` is now a numpy array of 100 values from 0 to 1. The derivative is now evaluated at 100 places with a single call.

△ **Jupyter Notebook Exercise 6.4:** Define $f(x) = x^3$. Use your `center` function to evaluate it's derivative $f'(x)$ in the x range $[-2, 2]$. Plot both $f(x)$ and $f'(x)$ in that range (in the same plot) using different colors for each. Add a legend and axis labels.

△ **Jupyter Notebook Exercise 6.5:** Define $f(x) = \sin(x)$ using the `np.sin` function. Use your `center` function to evaluate it's derivative $f'(x)$ in the x range $[0, 2\pi]$. Plot $f(x)$, $f'(x)$, and $\cos(x)$ in that range (all in the same plot) using different colors for each. Add a legend and axis labels.

6.4 Projectile Motion

In lecture, we derived the Euler Method for iteratively determining the trajectory of a particle:

$$\begin{aligned}\vec{v}_{n+1} &= \vec{v}_n + \tau \vec{a}_n \\ \vec{r}_{n+1} &= \vec{r}_n + \tau \vec{v}_n\end{aligned}$$

△ **Jupyter Notebook Exercise 6.6:** Implement a function

```
def euler(dt, x, y, vx, vy, ax, ay):
    # your code here
    return x, y, vx, vy
```

which calculates the next iteration of x, y, v_x , and v_y from the current values of x, y, v_x, v_y, a_x and a_y . Notice that this function returns several different variables at once using a comma separated list

referred to as tuple in python. To retrieve the individual variables from the tuple, simply call the function like this:

```
x,y,vx,vy = euler(x,y,vx,vy,ax,ay)
```

One downside of this convenient approach is that you must get the order of the variables correct! Check your implementation against the following test values:

```
print(np.around(euler(0.134, 0.659, 0.282, 0.662, 0.643, 0.900, 0.451),2))
print(np.around(euler(0.924, 0.959, 0.575, 0.299, 0.710, 0.699, 0.471),2))
```

and ensure that you get the correct output:

```
[0.75 0.37 0.78 0.7 ]
[1.24 1.23 0.94 1.15]
```

We will use the Euler Method to simulate projectile motion. We'll take the initial velocity to be 20 m/s and take $g = 9.8 \text{ m/s}^2$. Here's a snippet of code that sets these constants and determines the x and y coordinates of the initial velocity from an angle θ , which is set to 45° :

```
tau = 2*np.pi
vi   = 20    # [m/s]
g    = 9.8   # [m/s^2]
theta = tau/8
dt   = 0.01  # [s]
x    = 0     # [m]
y    = 0     # [m]
vx   = vi*np.cos(theta)
vy   = vi*np.sin(theta)
```

The trajectory of the particle can be determined using the following algorithm:

```
1   Create an empty array tjx # will contain x positions of the trajectory
2   Create an empty array tjy # will contain y positions of the trajectory
3   while y ≥ 0:
4       Append the x position to tjx
5       Append the y position to tjy
6       Compute the next values of x,y,vx and vy using the Euler Method
7   Plot tjy versus tjx
```

Notice that the algorithm stops just before the projectile reaches $y \leq 0$.

△ **Jupyter Notebook Exercise 6.7:** Use the Euler Method to plot the trajectory of a projectile with the initial conditions described above.

△ **Jupyter Notebook Exercise 6.8:** Derive an expression (paper and pencil) for the maximum range of the trajectory and evaluate the range for these initial conditions. Are the results consistent?

△ **Jupyter Notebook Exercise 6.9:** Extend your simulation to record v_x and v_y at each step along with the x and y positions. Take the mass of the projectile to be $m = 0.145 \text{ kg}$ and plot the kinetic energy, potential energy, and total energy as a function of time. To build an array containing the time of each step, for plotting quantities versus time, you can do:

```
t = np.arange(tjx.size)*dt
```

Include a legend. The x and y axes have changed to time and energy, so make certain to change the axes labels!

6.5 Projectile Motion with Drag

In this section we will consider the effect of air resistance on a baseball thrown at 20 m/s. We can model drag as a deceleration:

$$\vec{a} = -k|\vec{v}|\vec{v}$$

where $k = 0.00622 \text{ m}^{-1}$ for typical baseballs.

△ **Jupyter Notebook Exercise 6.10:** Extend your simulation to include the effect of drag. Plot the trajectory without drag and the trajectory with drag in the same plot. Include a legend and (as always) label all axes.

△ **Jupyter Notebook Exercise 6.11:** Including the effect of air resistance, plot the kinetic energy, potential energy, and total energy (kinetic plus potential) of the projectile as a function of time. Is the total energy of the projectile conserved?

Chapter 7

Pendulum Motion

7.1 Introduction

In this lab we will apply the Euler and Verlet methods to the pendulum problem. We will compare the results of the Verlet problem to the small angle approximation.

7.2 Preparation

Suppose you want to produce a plot of $f(t) = A \exp(-t/\lambda)$ versus t from $t = 0$ to 10 and $A = \lambda = 5$. For a visually smooth plot, you want about $N = 100$ points, but we'll start with a smaller number $N = 11$ for easy debugging. You create an array containing the 11 time values you want to plot:

```
MAX = 10
N = 11
t = np.linspace(0, MAX, N)
print(t)
```

You calculate the $y = f(t)$ values like this:

```
A = 5
LAMB = 5
y = A*np.exp(-t/LAMB) # !!!
print(y)
```

Make sure that you thoroughly understand the line marked “!!!”. That is calculating one y value for every t value, and so y is an array with the same shape and size as t :

```
print("t shape: ", np.shape(t), "y shape: ", np.shape(y))
print("t size: ", np.size(t), "y size: ", np.size(y))
```

With two arrays of the same shape, plotting them is a simple matter. Here we use the red line format, and add some axes labels:

```
plt.plot(t, y, "r-")
plt.xlabel("t (s)")
plt.ylabel("y")
```

If you look closely, you'll see kinks in the plot for $N = 11$. Increase to $N = 100$ for a visually smooth plot.

You are expected to do all of this on your own, from a prompt like this:

△ **Jupyter Notebook Exercise 7.1:** Plot the function $f(t) = A \exp(-t/\lambda)$ from $t = 0$ to 10, $A = 5$ and $\lambda = 5$ as a red line.

Now suppose instead that you already have a set of y -values in a `np.array` `yarr`, and you would like to plot y vs t , knowing that these y values were sampled starting at $t = 0$ with a uniform step size of $\tau = 0.2$ between each sample, i.e. at $t = 0, 0.2, 0.4, \dots$. In this case, you would create a `np.array` containing your time values as:

```
tau = 0.2
t = np.arange(yarr.size)*tau
```

△ **Jupyter Notebook Exercise 7.2:** Starting from the y values contained in:

```
yarr = np.array([1,1,1,2,4,7,3,2,1,1,0,0])
```

which you know correspond to time values starting at $t = 0$ with constant step size $\tau = 0.5$. Plot y vs t as blue circles and add axes labels.

7.3 Pendulum Motion

In lecture we showed a pendulum of length L can be described by the angle θ with respect to vertical (rest position), the angular velocity:

$$\omega = \frac{d\theta}{dt}$$

and the angular acceleration:

$$\alpha = \frac{d\omega}{dt} = \frac{d^2\theta}{dt^2}$$

In a constant gravitational field with acceleration g , the angular acceleration is:

$$\alpha = -\frac{g}{L} \sin \theta$$

Which we can write as:

$$\alpha = -\omega_L^2 \sin \theta$$

where

$$\omega_L = \sqrt{\frac{g}{L}}.$$

This problem, which is quite simple to pose, has no analytic solution in terms of elementary functions. Instead, we will rely on numerical techniques to solve it.

7.4 Small Angle Approximation

Analytic solutions are the precious gems that we use to validate our numerical techniques. We'll start our analysis in a region where we can solve the problem analytically. For small displacements of the pendulum, θ is small, and so:

$$\sin \theta \approx \theta$$




Figure 7.1: The small angle approximation for pendulum motion with period $T_L = 1$ s

and

$$\alpha = -\omega_L^2 \sin \theta \approx -\omega_L^2 \theta$$

The differential equation which describes the motion is:

$$\frac{d^2\theta}{dt^2} = -\omega_L^2 \theta$$

We showed in lecture that if the initial angular velocity is zero ($\omega_0 = 0$) and the initial theta position is θ_0 , then the solution is:

$$\theta(t) = \theta_0 \cos(\omega_L t) \quad (7.1)$$

as plotted in Fig. 7.1. The pendulum rocks back and forth following a sine wave with period:

$$T_L = \frac{2\pi}{\omega_L} = 2\pi\sqrt{\frac{L}{g}}$$

△ Jupyter Notebook Exercise 7.3: Calculate ω_L for a pendulum with $L = 1$ m and $g = 9.8$ m/s². What are the units of ω_L ?

△ Jupyter Notebook Exercise 7.4: Suppose you want to construct a pendulum such that for small oscillations the period $T_L = 1$ s. What should be the value of ω_L ? What length L should you use, assuming that $g = 9.8$ m/s?

△ **Jupyter Notebook Exercise 7.5:** Reproduce the plot in Fig. 7.1. Set the constants T_L , ω_L , θ_0 , and N as follows:

```
TL = 1          # period in seconds
wL = 2*np.pi/TL # small-angle angular frequency of pendulum
A = 0.2         # theta at t=0 (amplitude)
```

You must complete this problem **without using an explicit for loop** and only **six or less additional lines of code**. (Debugging print statements do not count toward this limit, use as many of those as you like!) To set the y axis label to the fancy θ do:

```
plt.ylabel("$\\theta$")
```

.

7.5 The Failure of the Euler Method

The Euler equations for angle θ and angular velocity ω are:

$$\theta_{n+1} = \theta_n + \tau \omega \quad (7.2)$$

$$\omega_{n+1} = \omega_n + \tau \alpha \quad (7.3)$$

where α is the angular acceleration and τ is the time step.

△ **Jupyter Notebook Exercise 7.6:** Implement a function

```
def euler(tau, theta, omega, alpha):
    # your code here
    return theta, omega
```

which implements one iteration of the Euler method. Test your `euler` function with these test values:

```
print(np.around(euler(-0.01, -0.28, -0.30, -0.06),3))
print(np.around(euler( 0.94,  0.32, -0.85, -0.86),3))
print(np.around(euler( 0.92, -0.16,  0.38, -0.32),3))
print(np.around(euler( 0.31,  0.12, -0.91, -0.76),3))
print(np.around(euler(-0.14,  0.96,  0.66, -0.73),3))
```

which should produce the following output:

```
[-0.277 -0.299]
[-0.479 -1.658]
[0.19  0.086]
[-0.162 -1.146]
[0.868 0.762]
```

You'll use this now thoroughly tested function more below. Don't change it!

△ **Jupyter Notebook Exercise 7.7:** Apply the Euler method to the problem of small oscillations of a pendulum with $T_L = 1$ s as in Fig. 7.1. First, set the parameters of your code just as before:

```
TL = 1          # period in seconds
wL = 2*np.pi/TL # small-angle angular frequency of pendulum
A = 0.2         # theta at t=0 (amplitude)
```

Then implement the Euler method as follows:

```

1   $\theta := A$ 
2   $\omega := 0$ 
3   $\tau := 0.0003$ 
4  Create an empty array tjth which will contain  $\theta$  positions of the trajectory
5  Append  $\theta$  to the array tjth.
6  Repeat  $N$  times:
7      Calculate  $\alpha = \omega_L^2 \theta$ 
8      Update  $\theta$  and  $\omega$  for acceleration  $\alpha$  and time step  $\tau$  by calling euler().
9      Append  $\theta$  to the array tjth.
10 Create an array t containing  $N+1$  appropriately spaced time values.
11 Plot tjth versus t

```

As always, first debug your code using a small value for N . Then, you should reproduce something that closely resembles Fig. 7.1 with $N = 10000$.

△ Jupyter Notebook Exercise 7.8: Repeat the exercise above (you can use cut and paste) with $\tau = 0.01$ and $N = 1000$. Yikes! Is energy conserved?

7.6 The Verlet Method

The Verlet Equation for this problem is:

$$\theta_{n+1} = 2\theta_n - \theta_{n-1} + \tau^2 \alpha \quad (7.4)$$

Notice that with the Verlet method, we will not need to calculate angular velocity ω in order to get the θ trajectory.

△ Jupyter Notebook Exercise 7.9: Implement a function

```

def verlet(tau, theta, oldth, alpha):
    # your code here
    return theta

```

which returns θ_{n+1} from $\theta_n = \text{theta}$ and $\theta_{n-1} = \text{oldth}$ using the verlet method. Test your `verlet` function with these test values:

```

print(np.around(verlet(-0.01, -0.28, -0.30, -0.06),3))
print(np.around(verlet( 0.94,  0.32, -0.85, -0.86),3))
print(np.around(verlet( 0.92, -0.16,  0.38, -0.32),3))
print(np.around(verlet( 0.31,  0.12, -0.91, -0.76),3))
print(np.around(verlet(-0.14,  0.96,  0.66, -0.73),3))

```

which should produce the following output:

```

-0.26
0.73
-0.971
1.077
1.246

```

You'll use this now thoroughly tested function more below. Don't change it!

△ **Jupyter Notebook Exercise 7.10:** In a previous exercise you showed that the Euler method, when applied to the problem of small oscillations of a pendulum with $T_L = 1$ s for $\tau = 0.01$ and $N = 1000$, is unstable. Instead, apply the Verlet method. You can reuse (by copying and pasting) much of your code from that previous exercise with a few changes:

- You will call your `verlet` function instead of `euler`.
- You no longer need to keep track of ω (`omega`)
- You will now have to keep track of two θ values at all times. At each update:

$$(\theta_n, \theta_{n-1}) \rightarrow (\theta_{n+1}, \theta_n)$$

- You can start things off with `oldth = theta = A`

With this method, you should produce many oscillations with no sign of instability.

△ **Jupyter Notebook Exercise 7.11:** So far, we have been using the small angle approximation. Modify your code to use the exact formula for the angular acceleration $\alpha = \omega_L^2 \sin \theta$ and set the initial position to $A = 2$. This is a trivial change to your numerical simulation, but it makes an analytic solution impossible!

△ **Jupyter Notebook Exercise 7.12:** (Optional Challenge) Run your Verlet analysis for $N = 1000$ steps for $A = 3$ and set τ appropriately so that you see a bit more than one period of motion. Plot the trajectory as a black line and read off the period T . Superimpose a plot of a cosine with amplitude A and period T .

△ **Jupyter Notebook Exercise 7.13:** (Optional Challenge) We've been starting things off with approximately zero angular velocity by setting `oldth = theta = A`. You can add velocity to the initial state with:

```
V = 1
theta = A
oldth = A - tau*V
```

Give the pendulum enough of a whack that it reaches all the way to top and keeps going. When plotting this trajectory, you can use `np.mod` to keep the `theta` values in the range from $-\pi$ to π if you want.

△ **Jupyter Notebook Exercise 7.14:** (Optional Challenge) Fix the Euler method for this pendulum problem by forcing energy to be conserved at each step. Compare your results to the Verlet method.

Chapter 8

Orbital Motion

8.1 Introduction

8.2 Orbital Motion

Chapter 9

Integration

9.1 Introduction

9.2 Integration

9.3 Gaussian Distribution

9.4 Elliptic Function